

Miscellaneous Thingies



Three Miscellaneous Things: Not on the final

- GPUs (Guest from Apple)
- Building a software pipeline in a 61C style
 - Or "How Nick pegged all cores on The Beast"
- Two cool ISA items
 - ARM Pointer Authentication Codes
 - RISC-V 16b instruction encoding

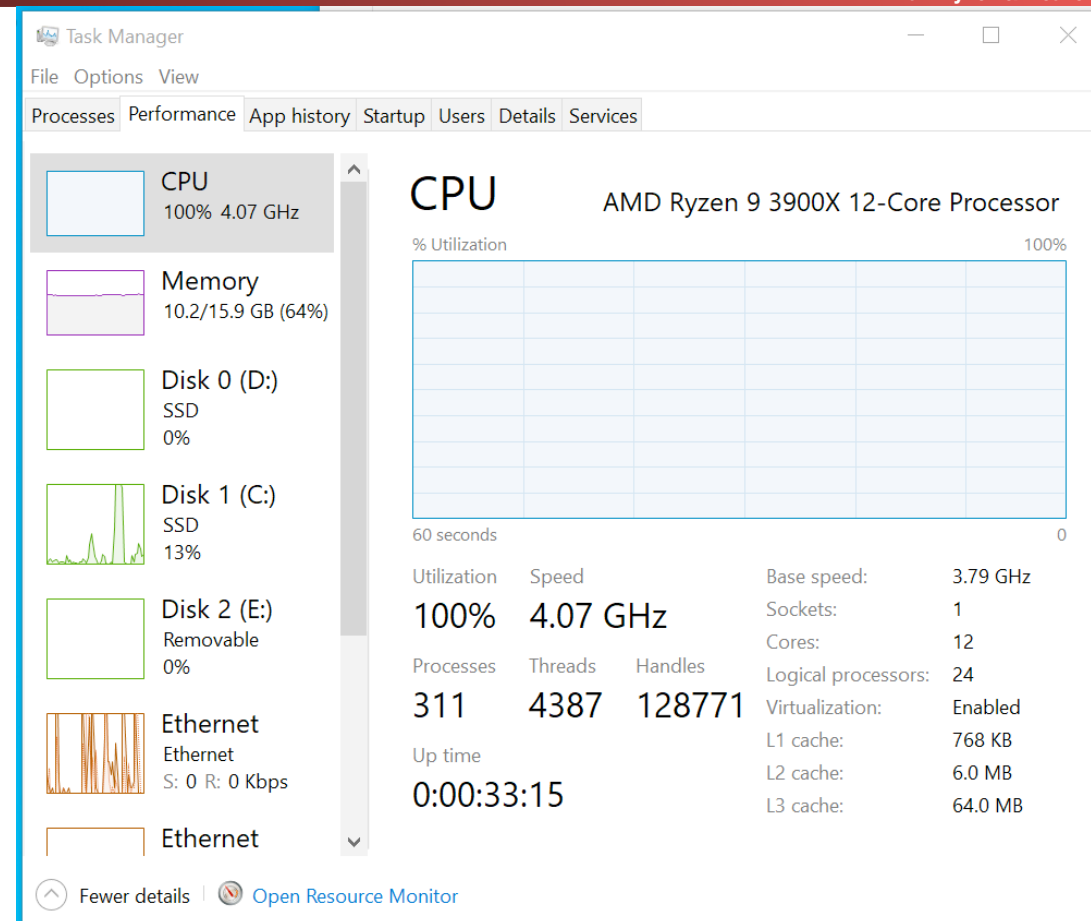
Guest Lecture Here... But Why?

- GPUs are huge SIMD parallel devices
 - "For every element do X" operations
- Really good for matrix multiplies
 - Which is what a huge fraction of "Machine Learning" really is
 - And how you render scenes graphically

Building A Data Pipeline...

61C Style on my Linux box

- Some bad actor stole >4 GB of data from UCOP
- Basically everything that was on the "secure" file transfer server in December
- The bad actor released at least some of this data **publicly**
 - As a 4GB compressed archive that anyone can download
- What information about **me** was in the archive?



What I need to know...

- I already know my social security # got breached
 - They told us that...
 - But I've got fraud alerts & freezes in place already
- But what other information?
 - Address? Phone #? Things I don't know about?
 - Tax information?
 - ***Banking information?***
 - The numbers on the bottom of a check are all an attacker needs to make fake checks

The Nature of the Dump...

- A **lot** of pdf files
 - PDFs are a pain to search, need to convert to text
- A **lot** of data tables
 - Some as comma-delimited text, some as excel spreadsheets, some in stada format
- Need to convert it to something reasonable
- Google around...
 - Nice linux OCR pipeline cobbled together:
PDF -> images -> OCR text
 - pandas can read both xlsx and stada files

Step 1:

File Conversion

- Want to convert everything into text files
- Obscenely parallel problem:
 - For every PDF do X
- But with some gotchas...
 - I can't just spawn 700 PDF->txt conversion programs
 - That would grind my machine to a halt by exceeding my available memory
 - And different invocations take a different amount of time
- Two approaches
 - Dynamically tune based on load...
 - When load < 95%, spawn another job and wait 1 second before checking again
 - When job ends and load < 99%, spawn another job
 - When job ends and load \geq 99%, don't spawn another job
 - Or just say "F-it, and keep X jobs live"

Keeping X jobs live: Fork/join with a limiter

- Used a simple go lang hack
 - Go == C with safety, garbage collection, and really good parallelism
 - Channels are like Queues in python:
You can add elements, and trying to read blocks until an element is available
- `capacity := make(chan bool, 10)`
`done := make(chan bool)`
- ```
func run(txt string) {
 c <- True; // Will block until capacity is available

 <- c; // will release capacity
 done <- True; // join on fork/join
}
```
- main just calls "go run()" on every line of stdin...  
and then an equal number of lines of "<- done"



# A bit of tuning...

- ~10 jobs pegs all CPU cores on the OCR pipeline...
  - And it took a couple hours to PDF->txt the lot:  
Driver used multiple threads for single documents
- The .xlsx and stata conversion was a lot faster...
  - Set to ~25 jobs instead of 10 (since python doesn't thread, especially on this task)
  - Took ~10 minutes or so
- Not fully efficient...
  - At the end of the PDF run there was no longer pegged CPU
  - 100% CPU utilization means efficiency loss due to context switching:  
Optimum would be ~95%
  - Also, really stressing the Windows virtualization...
  - I do all my work in "linux" under WSL

# And Now To Search

- Just use the same pipeline with grep...
- But...
  - Ends up **not** pegging the CPUs...  
Instead I'm pegging the "disk"!
- OK for just searching for me, but...
  - Want to be able to do a "for anyone who wants" service
- So to do this, parallelize on an alternate axis:
  - Don't check one person at a time, check **all** people using a single program
  - And then invoke that in parallel across all files
- Gotcha problem: Need to make sure to synchronize writes well
  - Again, go lang FTW:  
A channel for each user's results, the search does an atomic write to the channel

# Memory Hardening...

## ARM Pointer Authentication Codes

- Attackers want to overwrite memory...
  - When your C code fails to check a buffer


- The classic vulnerability

```
void foo() {
 char c[32];
 gets(c)
}
```

- Attacker gives you too long an input...
  - And c is stored on the stack
  - So the attacker overwrites not just c but the other stuff on the stack...
    - Such as the saved `ra`
  - The saved `ra` is overwritten to point to the attacker's code in memory

# Stack Canaries...

- Goal is to protect the return pointer from being overwritten by a stack buffer...
- When the program starts up, create a **random** value
  - The “stack canary”
- When starting a function, write it just below the saved frame pointer or ra register
- When returning in a function
  - First check the canary against the stored value



# Stack Canary Overhead...

- May require enabling an optional compiler flag...
  - So of course it is commonly not done!
- Requires a memory load & store on every function entrance
  - Highly cacheable so basically only 4 instructions on a typical RISC:
    - Load address of canary (2 instructions)
    - Load canary value into register
    - Store canary value onto stack
- Requires 2 memory loads and a (probably) not taken branch on exit
  - So 5 instructions on a typical RISC:
    - Load address
    - Load canary value
    - Load canary off stack
    - BNE (mark as probably-not-taken if you can)

# So example code...

- `la t0 canary # Reminder, turns into two`  
    `# instructions`  
    `lw t0 0(t0)`  
    `sw t0 x(sp) # four below where ra got stored`  
    `# if we don't bother saving the frame pointer`  
    `# if we do save the frame pointer, 4 below fp`
- `la t0 canary`  
    `lw t0 0(t0)`  
    `lw t1 x(sp)`  
    `bne t0 t1 dead_canary`  
        `# Make sure this is a forward branch:`  
        `# So CPU assumes it won't be taken`
- Note also generally sequential:  
    only parallelism present is in loading the canary from both the stack and storage

# Brute Force...

- Brute force: just simply try every possibility
  - Or if its a different random # each time, just always try the same number
- Even the smallest timeout goes along way:
  - If you can try 10,000 per second, trying  $2^{20}$  possibilities takes less than 2 minutes
  - If you can only try 10 per second, it takes a day and a half
  - And if 10 failures causes a 10 minute timeout...  
Forgettaboutit!
- Exponentials matter
  - If it take 1 minute to try  $2^{20}$ , it will take 16 hours to try  $2^{30}$
  - And 2 years to try  $2^{40}$ !
  - EG, Apple added a mitigation in the latest iOS:  
Crashing programs can (optionally) have an exponentially growing delay on restarting from crashes, which prevents attacks that need to repeatedly crash the service to extract information or get lucky

# Pointer Protection:

## Modern 64b ARM 8.3 Pointer Authentication

- <https://www.qualcomm.com/media/documents/files/whitepaper-pointer-authentication-on-armv8-3.pdf>
- ARM64 uses 64b pointers
- Idea: Since our pointers are 64b but we are only using say 42b of them...
  - Lets use that upper 22b to encrypt/protect pointers of various types!
- New instructions:
  - **PAC** -> Set Pointer Authentication Code
    - Sets the upper bits with a cryptographic checksum
  - **AUT** -> Check and Remove Pointer Authentication Code
    - If the check is invalid, it will instead put an error in the checksum space:  
If the pointer is dereferenced it causes an error
  - **XAUT** -> Strip PAC without checking
- Instructions are in NO-OP space if the processor doesn't support them



# Plus some non-NOOP higher performance options

- When you know you will be running on a processor which supports it
  - check & return:  
Check the return address has a valid PAC and if so, return
  - check & load:  
Check the PAC and if so, load the pointer
  - check & branch:  
Check the PAC and if so, do a jump-and-link to that pointer
- Allows the complete elimination of the overhead for checking!
  - Well, cheat: You cause it to trigger an exception on instruction committing and just assume the pointer is valid to start with...

# How To Use...

- There are 5 secrets for pointer protection
  - These contain random 128b secrets that are used to authenticate the pointer:  
Provided by the OS
  - Two for data (DA/DB), two for instruction (IA/IB), and one general purpose (GA)
- These are contained in processor registers (ARM equivalent of the CSRs),  
and are ***not readable to the program itself!***
  - Key property: An information leakage vulnerability can't defeat this protection on a user-level program
  - But it could on a kernel level program:  
Solution would be to also have a secret random to the CPU that is included but non readable
- Other workaround: find a vulnerability that can trick the program into authenticating new pointers it shouldn't, or be able to reuse authenticated pointers in another context

# So in practice

- The PAC is a function of the pointer, an additional register (or register 0) and the hidden secret
  - **PACIA x30 sp**  
**AUTIA x30 sp**  
Protect/Authenticate x30 as a function of x30, sp, and the secret data associated with the Instruction A context (x30 is the default link register for ARM == **ra** in RISC-V)
- Thanks to crypto-magic discussed in 161, the PAC's "look random"
  - Changing a single bit of anything should result in something looking totally different and random
- So to guess a 22 bit PAC would be 1 in 4 million odds.

# So Cheaper Stack Canaries...

- On function entry: Create the PAC for the return address
  - Using the stack pointer as the context itself:  
This means the return address can't even be moved
- On function exit: Check & return as normal
- With backwards compatibility: only 2 instructions
  - **PACIA** on function start, **AUTIA** on function end
- Without backwards compatibility: only 1 instruction!
  - Just the **PACIA** on function start and a check & return on exit
  - Saves 8 instructions... Or >85%!
- Only 22 bits of entropy but...
  - If you get more than a few failures, just keep the program dead!

# Or Protecting vtable pointers...

- When you allocate a new C++ object...
  - The first thing is a "vtable pointer", really a pointer to an array of pointers to functions
    - Attackers want to overwrite this with their own version
  - Protect the vtable pointer with a context and register 0:  
One additional instruction when calling `new()`
  - Then have the vtable itself live in read-only space so it can't be overwritten
- Now when calling a virtual function...
  - Check & Load the vtable pointer (RISC-V like pseudocode):  
eg, if the object pointer is in `s0`, the vtable pointer is at the start of `s0`...  
`LDRAA t0 0(s0) # Load 0 + s0, authenticated with data A`  
`LW t0 X(t0) # X == the specific function to call`  
`JALR t0 # Actually call it`
- Now you **can't** overwrite a C++ object's vtable pointer to something else without either being very lucky, finding a separate vulnerability, or replacing with another valid pointer that you acquire...  
And the overhead is literally **nothing**!
  - Apart from you need to recompile and using the latest ARM silicon, that is

# Probably the biggest benefit for Apple going to ARM

- MacOSX ARM will be able to ***assume*** PAC support!
  - Since it is Apple A12 or newer processors only
  - Latest iOS also just started really aggressively turning on PAC support
- Can therefore use the more efficient primitives:
  - Check & Branch Register, Check & Load, Check & Return  
which all eliminate the instruction needed in a separate check
  - Usable in both the kernel and user space:  
Acts to harden both applications and the underlying OS
- x86 has nothing like this in the pipeline!
- If you have a choice of architecture for a product: ARM 8.3+
  - This gives you so much real-world security for crappy C-code

# RISC-V 16b ISA

- Observation:  
Although we encode instructions with 32b, a lot of the instructions follow common patterns
  - Some registers used a lot more than others
  - Immediates are often small
  - Same source and destination for 3-operand operations
- So the optional "C" instruction set can be mixed in on a per-instruction basis
  - Look at the first two bits of an instruction can determine its type
    - 32b or 16b
  - For a C instruction,  $PC \leftarrow PC + 2$  instead of  $PC + 4$ 
    - And now normal instructions only need to be half-word aligned
- Results in ~30% smaller code
  - Roughly the same performance gain as **doubling** the icache!

# The Instruction Encoding

| Format | Meaning              | 15     | 14 | 13          | 12 | 11     | 10   | 9 | 8 | 7      | 6 | 5    | 4 | 3  | 2 | 1  | 0 |  |  |
|--------|----------------------|--------|----|-------------|----|--------|------|---|---|--------|---|------|---|----|---|----|---|--|--|
| CR     | Register             | funct4 |    |             |    | rd/rs1 |      |   |   | rs2    |   |      |   | op |   |    |   |  |  |
| CI     | Immediate            | funct3 |    | imm         |    | rd/rs1 |      |   |   | imm    |   |      |   | op |   |    |   |  |  |
| CSS    | Stack-relative Store | funct3 |    | imm         |    |        |      |   |   | rs2    |   |      |   | op |   |    |   |  |  |
| CIW    | Wide Immediate       | funct3 |    | imm         |    |        |      |   |   |        |   | rd'  |   | op |   |    |   |  |  |
| CL     | Load                 | funct3 |    | imm         |    |        | rs1' |   |   | imm    |   | rd'  |   | op |   |    |   |  |  |
| CS     | Store                | funct3 |    | imm         |    |        | rs1' |   |   | imm    |   | rs2' |   | op |   |    |   |  |  |
| CB     | Branch               | funct3 |    | offset      |    |        | rs1' |   |   | offset |   |      |   | op |   |    |   |  |  |
| CJ     | Jump                 | funct3 |    | jump target |    |        |      |   |   |        |   |      |   |    |   | op |   |  |  |

RVC Register Number

Integer Register Number

Integer Register ABI Name

Floating-Point Register Number

Floating-Point Register ABI Name

|     |     |     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|-----|-----|
| 000 | 001 | 010 | 011 | 100 | 101 | 110 | 111 |
| x8  | x9  | x10 | x11 | x12 | x13 | x14 | x15 |
| s0  | s1  | a0  | a1  | a2  | a3  | a4  | a5  |
| f8  | f9  | f10 | f11 | f12 | f13 | f14 | f15 |
| fs0 | fs1 | fa0 | fa1 | fa2 | fa3 | fa4 | fa5 |



# Stack Relative Load/Stores:

## Shrink preamble/postamble code by nearly 50%!

- Immediates are 0-extended
  - Because we write up from the stack pointer
- Immediates assume basic alignment (lower two bits 0 for words, three for doubles...)

| 15      | 13 | 12 | 11        | 7             | 6 | 2               | 1 | 0  |
|---------|----|----|-----------|---------------|---|-----------------|---|----|
| funct3  |    |    | imm       | rd            |   | imm             |   | op |
| 3       |    |    | 1         | 5             |   | 5               |   | 2  |
| C.LWSP  |    |    | offset[5] | dest $\neq$ 0 |   | offset[4:2 7:6] |   | C2 |
| C.LDSP  |    |    | offset[5] | dest $\neq$ 0 |   | offset[4:3 8:6] |   | C2 |
| C.LQSP  |    |    | offset[5] | dest $\neq$ 0 |   | offset[4 9:6]   |   | C2 |
| C.FLWSP |    |    | offset[5] | dest          |   | offset[4:2 7:6] |   | C2 |
| C.FLDSP |    |    | offset[5] | dest          |   | offset[4:3 8:6] |   | C2 |

| 15      | 13 | 12 | 7               | 6 | 2   | 1 | 0  |
|---------|----|----|-----------------|---|-----|---|----|
| funct3  |    |    | imm             |   | rs2 |   | op |
| 3       |    |    | 6               |   | 5   |   | 2  |
| C.SWSP  |    |    | offset[5:2 7:6] |   | src |   | C2 |
| C.SDSP  |    |    | offset[5:3 8:6] |   | src |   | C2 |
| C.SQSP  |    |    | offset[5:4 9:6] |   | src |   | C2 |
| C.FSWSP |    |    | offset[5:2 7:6] |   | src |   | C2 |
| C.FSDSP |    |    | offset[5:3 8:6] |   | src |   | C2 |

# Register Relative Loads & Stores

- Same zero-extending and alignment tricks

| 15     | 13 12         | 10 9 | 7 6         | 5 4  | 2 1 | 0 |
|--------|---------------|------|-------------|------|-----|---|
| funct3 | imm           | rs1' | imm         | rd'  | op  |   |
| 3      | 3             | 3    | 2           | 3    | 2   |   |
| C.LW   | offset[5:3]   | base | offset[2 6] | dest | C0  |   |
| C.LD   | offset[5:3]   | base | offset[7:6] | dest | C0  |   |
| C.LQ   | offset[5 4 8] | base | offset[7:6] | dest | C0  |   |
| C.FLW  | offset[5:3]   | base | offset[2 6] | dest | C0  |   |
| C.FLD  | offset[5:3]   | base | offset[7:6] | dest | C0  |   |

| 15     | 13 12         | 10 9 | 7 6         | 5 4  | 2 1 | 0 |
|--------|---------------|------|-------------|------|-----|---|
| funct3 | imm           | rs1' | imm         | rs2' | op  |   |
| 3      | 3             | 3    | 2           | 3    | 2   |   |
| C.SW   | offset[5:3]   | base | offset[2 6] | src  | C0  |   |
| C.SD   | offset[5:3]   | base | offset[7:6] | src  | C0  |   |
| C.SQ   | offset[5 4 8] | base | offset[7:6] | src  | C0  |   |
| C.FSW  | offset[5:3]   | base | offset[2 6] | src  | C0  |   |
| C.FSD  | offset[5:3]   | base | offset[7:6] | src  | C0  |   |

# Jumps & Branches

| 15     | 13 12                         | 2 1 | 0 |
|--------|-------------------------------|-----|---|
| funct3 | imm                           | op  |   |
| 3      | 11                            | 2   |   |
| C.J    | offset[11 4 9:8 10 6 7 3:1 5] | C1  |   |
| C.JAL  | offset[11 4 9:8 10 6 7 3:1 5] | C1  |   |

| 15     | 12 11 | 7 6 | 2 1 | 0 |
|--------|-------|-----|-----|---|
| funct4 | rs1   | rs2 | op  |   |
| 4      | 5     | 5   | 2   |   |
| C.JR   | src≠0 | 0   | C2  |   |
| C.JALR | src≠0 | 0   | C2  |   |

| 15     | 13 12         | 10 9 | 7 6               | 2 1 | 0 |
|--------|---------------|------|-------------------|-----|---|
| funct3 | imm           | rs1' | imm               | op  |   |
| 3      | 3             | 3    | 5                 | 2   |   |
| C.BEQZ | offset[8 4:3] | src  | offset[7:6 2:1 5] | C1  |   |
| C.BNEZ | offset[8 4:3] | src  | offset[7:6 2:1 5] | C1  |   |

# And Then Assorted ALU instructions...

- Load 6 bit immediate values
  - Either to the lower 6 bits or bits 17-12
    - Useful for smaller immediates
- ADDI to self with small immediate
  - And a special form for a scaled by 16 immediate to the stack pointer
  - And another one to add a 0-extended immediate to the stack pointer to get addresses of stack-allocated variables
- Left shift more important than right shift:
  - Can left shift any register, but right shift only the encoded 8...
- Two register operations rather than 3 for the basic ALU ops:
  - Add, subtract, and, or, xor

# Result is massive savings...

- Program preamble/postamble
  - Basically shrunk by 50%
- Common array offset access
  - How many times do you access `x[0]` or `x[small positive integer]`?
- Shorter jump offsets
- Branch = or != 0
  - So all `beq x0` and `bne x0...`
- Control logic complexity very low even for an aggressive processor
  - Easy decision tree to know length of instructions compared with something like x86